


```
>>> c++ deep-learning.cpp
```


```
>>> ./a.out data.csv
```

```
>>> 42
```




OSI Model of Machine Learning

GPUs




Teams using GPUs at ImageNet Competition

TITAN X FOR DEEP LEARNING




Training Time for ImageNet

GPUs

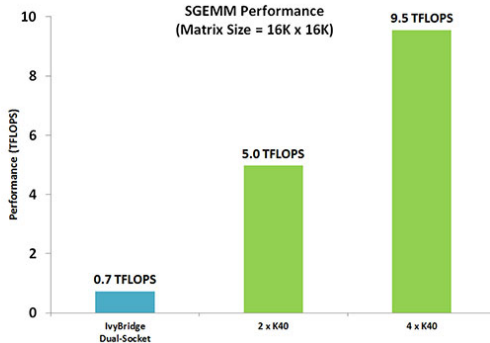


Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour



CPU vs. GPU Performance per Watt

GPUs



Single Precision General Matrix/Matrix Multiply Performance